


Lecture 5


Machine Learning Adv.

Jaeyun Kang


Overfitting



Underfit



Just Right




Overfit

Overfitting

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

If loss goes to 0,


trained model captures
all **training data**'s pattern



Overfitting

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

How about **new data**?




Overfitting

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

If the model is **over-fitting**
the pattern of training data,


it can't find **common data pattern**



Overfitting

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$


common data pattern



Overfitting

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

How can we train the model
to find **common data pattern**?




Overfitting - Regularization

$$L_i = -\log\left(\frac{e^{s y_i}}{\sum_j e^{s_j}}\right)$$

How can we train the model
to find **common data pattern**?

= **Regularization**




Overfitting - Regularization

$$L_i = -\log\left(\frac{e^{s y_i}}{\sum_j e^{s_j}}\right)$$

For **Regularization**,

The **variance** of **W** matrix
must be small



Interpreting a Linear Classifier (Review)



$$f(x, W) = Wx$$

Example trained weights of a
linear classifier trained on
CIFAR-10:

Interpreting a Linear Classifier (Review)



$$f(x, W) = Wx$$

**W matrix with small variance
can represent the
common data pattern**




Overfitting - Regularization

$$L_i = -\log\left(\frac{e^{s y_i}}{\sum_j e^{s_j}}\right)$$

For **Regularization**,

The **variance** of **W** matrix
must be small



Overfitting - Regularization

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

For **Regularization**.

The **variance** of **W** matrix
must be small

$$x = [1, 1, 1, 1]$$

$$w_1 = [4, 0, 0, 0] \quad ?$$

$$w_1^T x = w_2^T x = 4$$

$$w_2 = [1, 1, 1, 1] \quad ?$$

Overfitting - Regularization

$$L_i = -\log\left(\frac{e^{s y_i}}{\sum_j e^{s_j}}\right)$$

For **Regularization**.

The **variance** of **W** matrix
must be small

$$x = [1, 1, 1, 1]$$

$$w_1 = [4, 0, 0, 0]$$

$$w_1^T x = w_2^T x = 4$$

$$w_2 = [1, 1, 1, 1]$$

Overfitting - Regularization

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

How can we adjust
variance of **W** matrix ?

For **Regularization**.

The **variance** of **W** matrix
must be small

$$w_1 = [4, 0, 0, 0]$$

$$w_2 = [1, 1, 1, 1]$$

Overfitting - Regularization

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

For **Regularization**.

The **variance** of **W** matrix
must be small

How can we adjust
variance of W matrix ?

minimize **square sum** of
w components!

$$w_1 = [4, 0, 0, 0]$$

$$w_2 = [1, 1, 1, 1]$$

Overfitting - Regularization

$$L_i = -\log \left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}} \right) + \lambda \sum_k \sum_l w_{k,l}^2$$

How can we adjust
variance of W matrix ?
minimize **square sum** of
w components!

$$w_1 = [4, 0, 0, 0]$$

$$w_2 = [1, 1, 1, 1]$$

Overfitting - Regularization

$$L_i = -\log \left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}} \right) + \overset{\text{Regularization Strength}}{\lambda} \sum_k \sum_l W_{k,l}^2$$

How can we adjust
variance of W matrix ?
minimize **square sum** of
w components!

$$w_1 = [4, 0, 0, 0]$$

$$w_2 = [1, 1, 1, 1]$$

Evaluation

**How we evaluate the model's performance?
(How well trained?)**

Evaluation

How we evaluate the model's performance?
(How well trained?)

$\text{num}(\text{predicted_label} == \text{true_label}) / \text{total_num} * 100 (\%)$

Evaluation

How we evaluate the model's performance?
(How well trained?)

$\text{num}(\text{predicted_label} == \text{true_label}) / \text{total_num} * 100 (\%)$

But, not for **training dataset**

Evaluation


How we evaluate the model's performance?
(How well trained?)

$\text{num}(\text{predicted_label} == \text{true_label}) / \text{total_num} * 100 (\%)$



But, not for **training dataset**

Overfitting may have occurred !


Training Set and Test Set



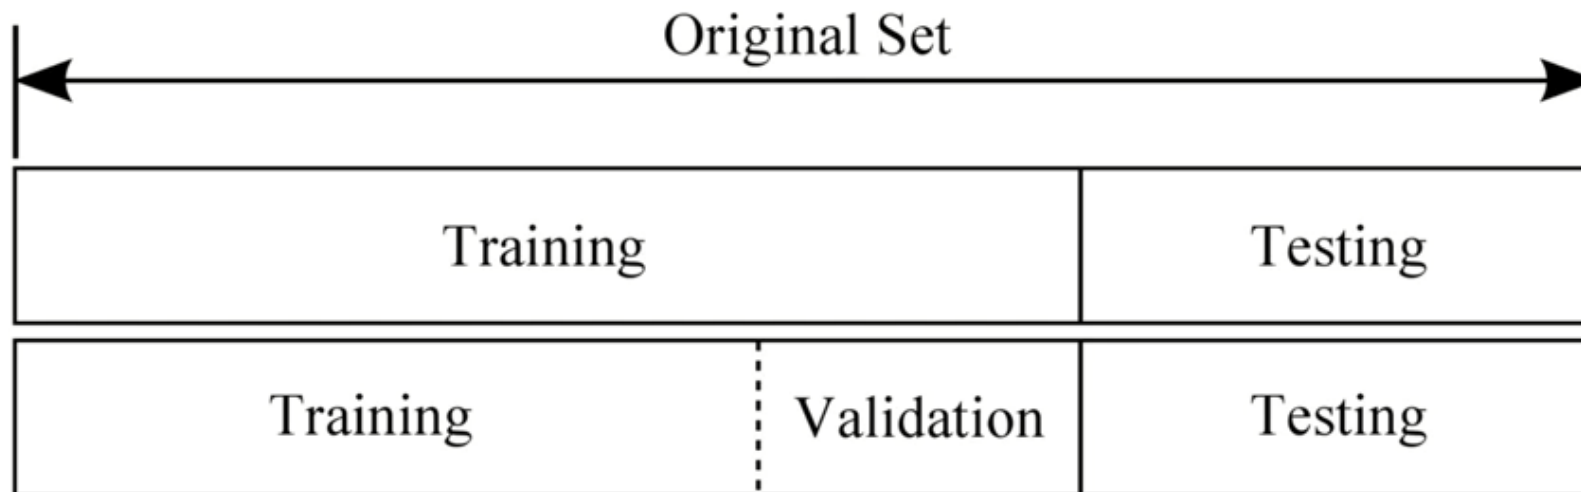
Split Dataset into
Training Set for Training
Test Set for Testing (Predicting)



Overfitting and Training-Test Error



Validation Set



α, λ

learning rate, regularization strength

Hyperparameter Tuning

Tensorflow Practice - MNIST

```
r = randint(0, mnist.test.num_examples - 1)
plt.imshow(mnist.test.images[r:r+1].reshape(28, 28), cmap='Greys', interpolation='nearest')
plt.show()
print("Prediction: ", sess.run(tf.argmax(scores, 1), feed_dict={x: mnist.test.images[r:r+1]}))
```

```
correct_prediction = tf.equal(tf.argmax(prob, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print(sess.run(accuracy, feed_dict={x: mnist.test.images, y: mnist.test.labels}))
```

Tensorflow Practice - MNIST

```
r = randint(0, mnist.test.num_examples - 1)
plt.imshow(mnist.test.images[r:r+1].reshape(28, 28), cmap='Greys', interpolation='nearest')
plt.show()
print("Prediction: ", sess.run(tf.argmax(scores, 1), feed_dict={x: mnist.test.images[r:r+1]}))
```

```
correct_prediction = tf.equal(tf.argmax(prob, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print(sess.run(accuracy, feed_dict={x: mnist.test.images, y: mnist.test.labels}))
```

loss: 0.280254434008

loss: 0.277931706214


loss: 0.277339565713

Prediction: [6]

0.9189

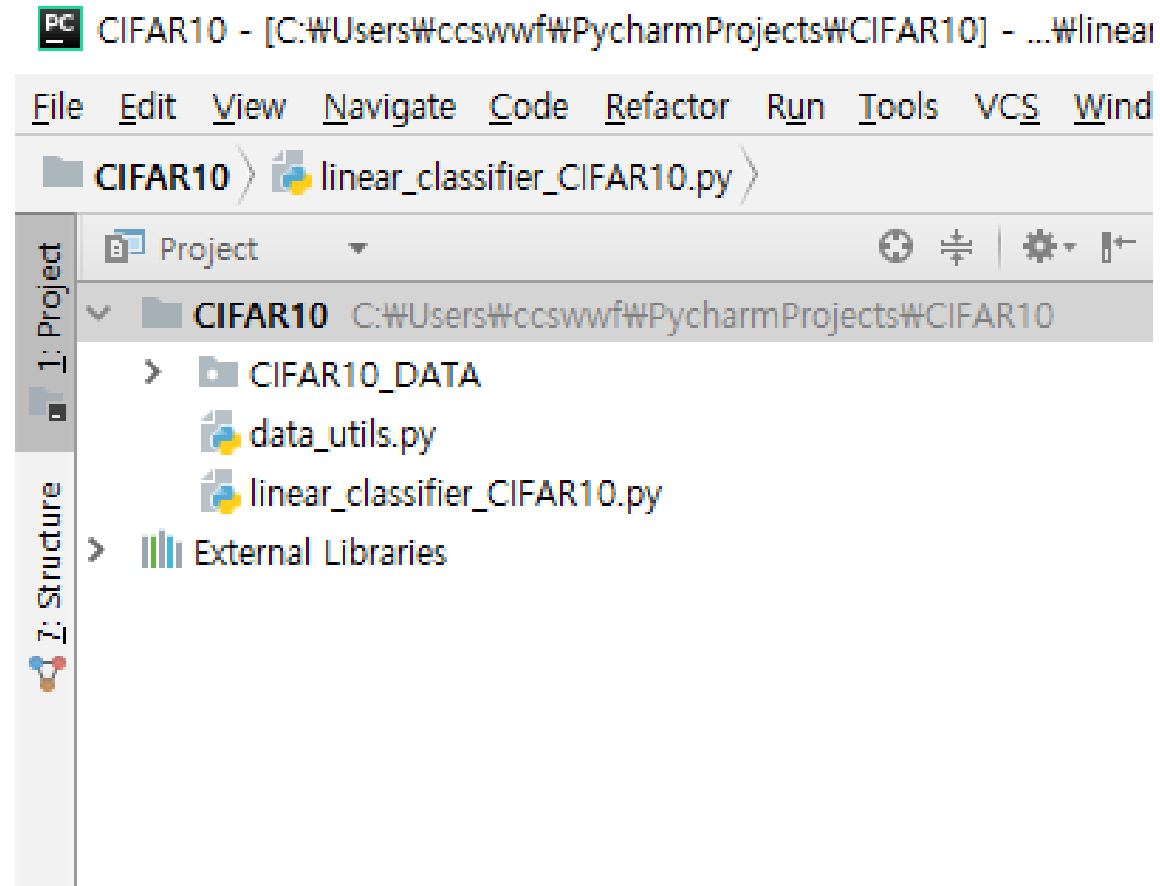
92%

Tensorflow Practice - CIFAR10



**60000 32x32 colour images
in 10 classes**

Tensorflow Practice - CIFAR10



Tensorflow Practice - CIFAR10

```
data_utils.py x linear_classifier_CIFAR10.py x
1 from data_utils import load_CIFAR10
2 import matplotlib.pyplot as plt
3 import numpy as np
4 import tensorflow as tf
5
6 # Load the raw CIFAR-10 data.
7 cifar10_dir = 'CIFAR10_DATA/cifar-10-batches-py'
8 X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)
9
10 # As a sanity check, we print out the size of the training and test data.
11 print('Training data shape: ', X_train.shape)
12 print('Training labels shape: ', y_train.shape)
13 print('Test data shape: ', X_test.shape)
14 print('Test labels shape: ', y_test.shape)
15 print('')
```

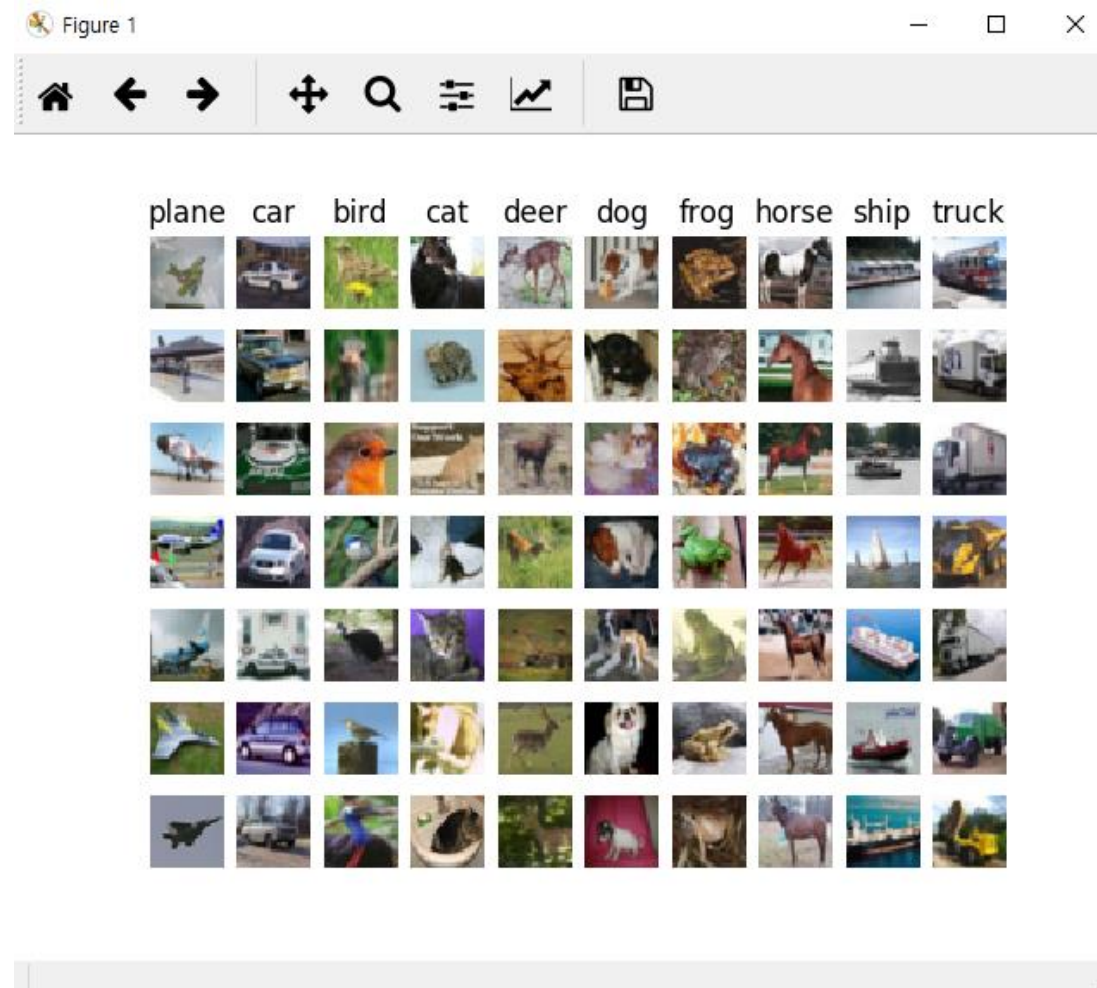
```
linear_classifier_CIFAR10
C:\Users\ccswwf\Anaconda3\python.exe C:/Use
Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
```

Tensorflow Practice - CIFAR10

```
# Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
num_classes = len(classes)
samples_per_class = 7

for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()
```

Tensorflow Practice - CIFAR10



Tensorflow Practice - CIFAR10

```
num_training = 49000
num_test = 1000

# Make training set
mask = range(num_training)
X_train = X_train[mask]
y_train = y_train[mask]

# Make test set
mask = range(num_test)
X_test = X_test[mask]
y_test = y_test[mask]
```

Tensorflow Practice - CIFAR10

```
num_training = 49000  
num_test = 1000
```

```
# Make training set  
mask = range(num_training)  
X_train = X_train[mask]  
y_train = y_train[mask]
```

```
# Make test set  
mask = range(num_test)  
X_test = X_test[mask]  
y_test = y_test[mask]
```

```
# Vectorize X matrix  
X_train = np.reshape(X_train, (X_train.shape[0], -1))  
X_test = np.reshape(X_test, (X_test.shape[0], -1))
```

Tensorflow Practice - CIFAR10

```
num_training = 49000  
num_test = 1000
```

```
# Make training set  
mask = range(num_training)  
X_train = X_train[mask]  
y_train = y_train[mask]
```

```
# Make test set  
mask = range(num_test)  
X_test = X_test[mask]  
y_test = y_test[mask]
```

```
# Vectorize X matrix  
X_train = np.reshape(X_train, (X_train.shape[0], -1))  
X_test = np.reshape(X_test, (X_test.shape[0], -1))
```

```
# One-hot encoding for train labels  
Y_train = np.zeros((y_train.shape[0], 10))  
Y_train[np.arange(y_train.shape[0]), y_train] = 1  
y_train = Y_train
```

```
# One-hot encoding for test labels  
Y_test = np.zeros((y_test.shape[0], 10))  
Y_test[np.arange(y_test.shape[0]), y_test] = 1  
y_test = Y_test
```

Tensorflow Practice - CIFAR10

```
num_training = 49000  
num_test = 1000
```

```
# Make training set  
mask = range(num_training)  
X_train = X_train[mask]  
y_train = y_train[mask]
```

```
# Make test set  
mask = range(num_test)  
X_test = X_test[mask]  
y_test = y_test[mask]
```

```
# Vectorize X matrix  
X_train = np.reshape(X_train, (X_train.shape[0], -1))  
X_test = np.reshape(X_test, (X_test.shape[0], -1))
```

```
# One-hot encoding for train labels  
Y_train = np.zeros((y_train.shape[0], 10))  
Y_train[np.arange(y_train.shape[0]), y_train] = 1  
y_train = Y_train
```

```
# One-hot encoding for test labels  
Y_test = np.zeros((y_test.shape[0], 10))  
Y_test[np.arange(y_test.shape[0]), y_test] = 1  
y_test = Y_test
```

Tensorflow Practice - CIFAR10

```
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

Tensorflow Practice - CIFAR10

```
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Train data shape: (49000, 3072)
```

```
Train labels shape: (49000, 10)
```

```
Test data shape: (1000, 3072)
```


```
Test labels shape: (1000, 10)
```

Tensorflow Practice - CIFAR10

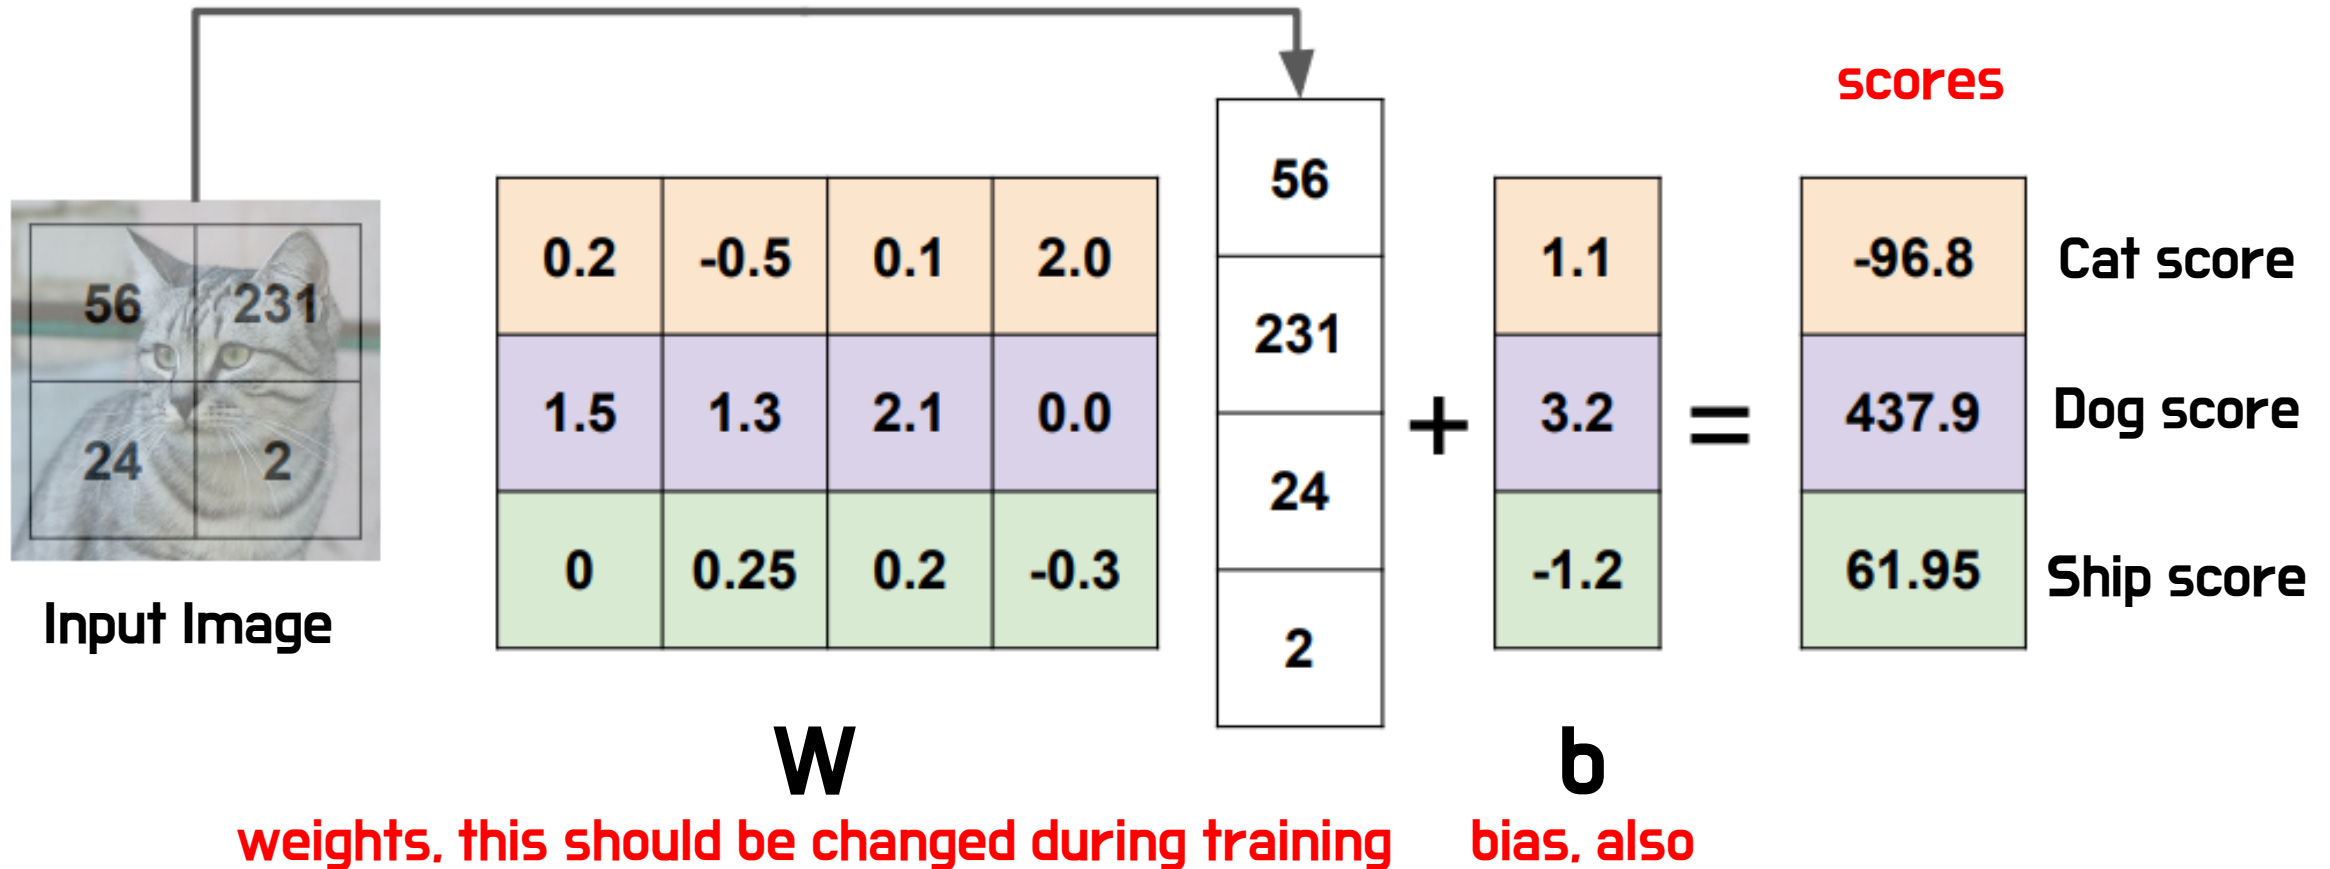
```
# Preprocessing: subtract the mean image (for zero-centered data)
# first: compute the image mean based on the training data
mean_image = np.mean(X_train, axis=0)
plt.figure(figsize=(4,4))
plt.imshow(mean_image.reshape((32,32,3)).astype('uint8')) # visualize the mean image
plt.show()

# second: subtract the mean image from train and test data
X_train -= mean_image
X_test -= mean_image
```

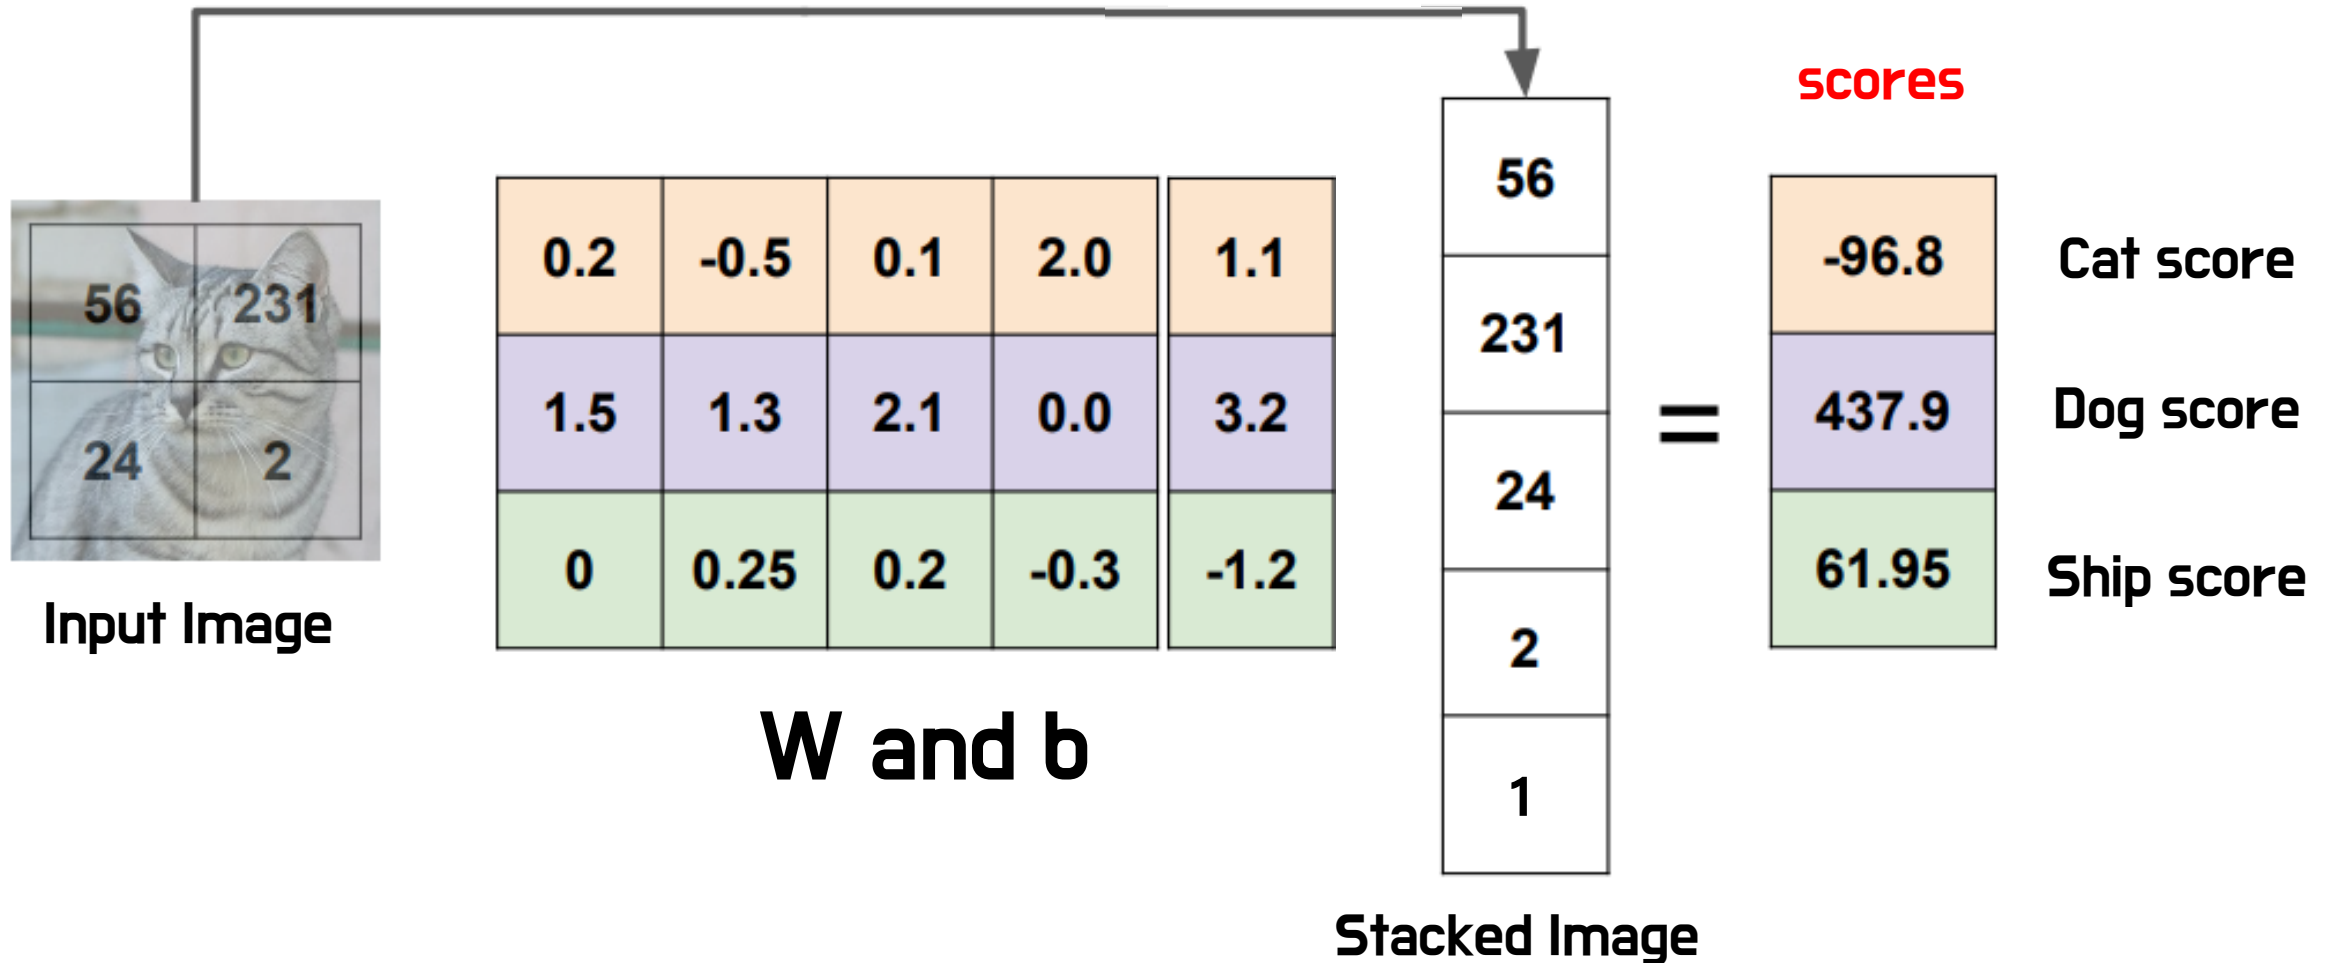
Tensorflow Practice - CIFAR10



Tensorflow Practice - CIFAR10



Tensorflow Practice - CIFAR10



Tensorflow Practice - CIFAR10

```
# third: append the bias dimension of ones (i.e. bias trick) so that our SVM  
# only has to worry about optimizing a single weight matrix  $W$ .  
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])  
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])  
  
print(X_train.shape, X_test.shape)
```

(49000, 3073) (1000, 3073)

Tensorflow Practice - CIFAR10

```
learning_rate = 1e-7
reg_strength = 5e4

x = tf.placeholder(tf.float32, [None, 3073])
y = tf.placeholder(tf.float32, [None, 10])

# Make weight matrix (includes 'b'(bias) vector)
W = tf.Variable(tf.random_normal(shape=(3073, 10), mean=0.0, stddev=0.001, dtype=tf.float32))

scores = tf.matmul(x, W)
prob = tf.nn.softmax(scores)

cross_entropy_loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels=y, logits=scores))
#cross_entropy_loss += reg_strength * tf.reduce_sum(tf.square(W))

train = tf.train.GradientDescentOptimizer(learning_rate).minimize(cross_entropy_loss)
```

Tensorflow Practice - CIFAR10

```
with tf.Session() as sess:
    sess.run(init)
    loss_hist = []
    for iter in range(1500):
        batch_size = 200
        rand_idx = np.random.choice(X_train.shape[0], batch_size)

        batch_x = X_train[rand_idx]
        batch_y = y_train[rand_idx]

        loss, _ = sess.run([cross_entropy_loss, train], feed_dict={x:batch_x, y:batch_y})
        loss_hist.append(loss)

    if iter % 100 == 0:
        print("iteration: ", iter, " loss: ", loss)
```

Tensorflow Practice - CIFAR10

```
iteration: 0  loss: 5.59391
iteration: 100  loss: 3.89708
iteration: 200  loss: 3.26234
iteration: 300  loss: 3.50787
iteration: 400  loss: 3.03151
iteration: 500  loss: 2.97595
iteration: 600  loss: 3.18339
iteration: 700  loss: 2.99337
```


Tensorflow Practice - CIFAR10

```
loss, _ = sess.run([cross_entropy_loss, train], feed_dict={x:batch_x, y:batch_y})
loss_hist.append(loss)

if iter % 100 == 0:
    print("iteration: ", iter, " loss: ", loss)

plt.plot(loss_hist)
plt.xlabel('Iteration number')
plt.ylabel('Loss value')
plt.show()
```

Tensorflow Practice - CIFAR10



Tensorflow Practice - CIFAR10

```
plt.show()
```

```
# Evaluate the model with test set
correct_prediction = tf.equal(tf.argmax(prob, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print(sess.run(accuracy, feed_dict={x: X_test, y: y_test}))
```

```
iteration: 1200 loss: 2.63387
iteration: 1300 loss: 2.62079
iteration: 1400 loss: 2.88722
0.267
```

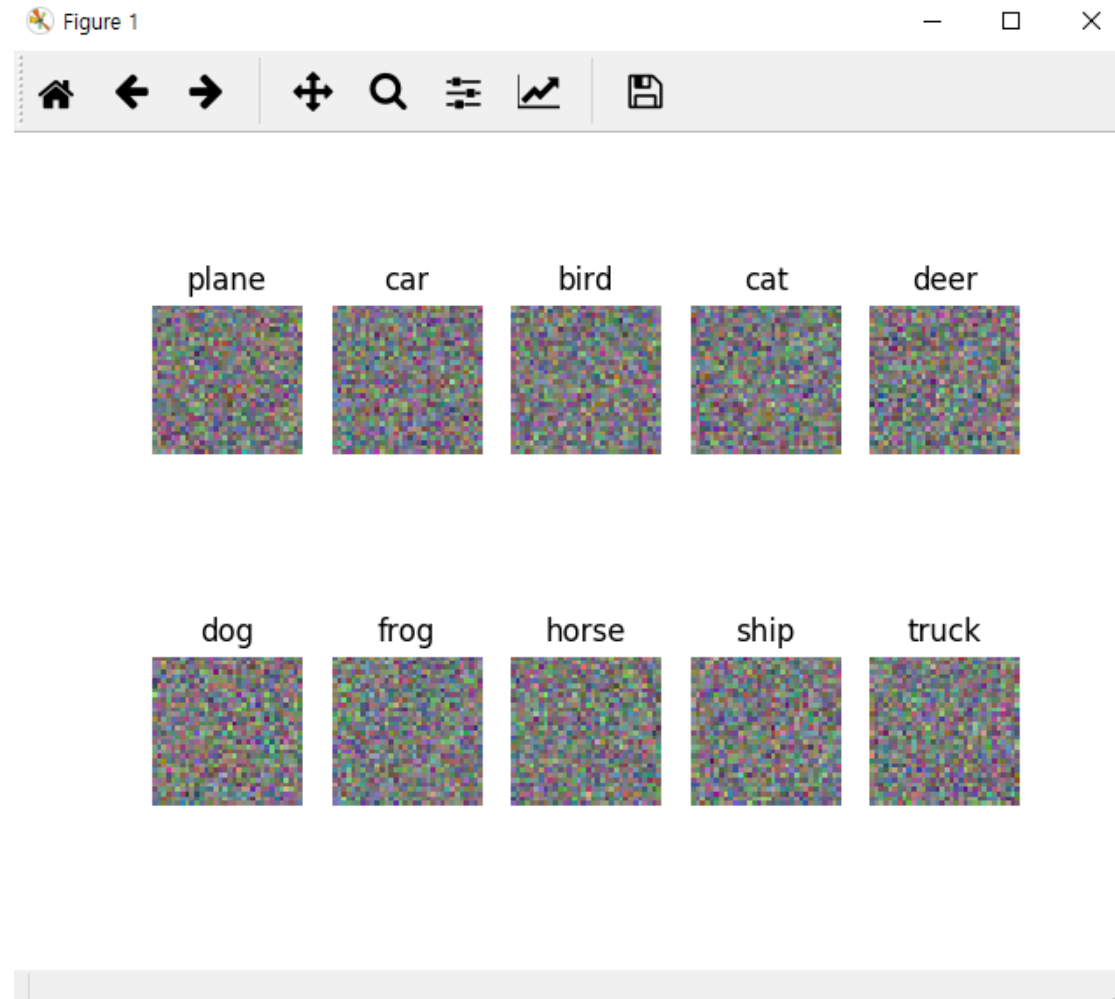
26.7%

Tensorflow Practice - CIFAR10

```
print(sess.run(accuracy, feed_dict={x: X_test, y: y_test}))

# Visualize the trained weight
w = sess.run(W)[:,-1, :] # strip out the bias
w = w.reshape(32, 32, 3, 10)
w_min, w_max = np.min(w), np.max(w)
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)
    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])
plt.show()
```

Tensorflow Practice - CIFAR10



Tensorflow Practice - CIFAR10


```
cross_entropy_loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels=y, logits=scores))  
cross_entropy_loss += reg_strength * tf.reduce_sum(tf.square(W))
```

Tensorflow Practice - CIFAR10


```
cross_entropy_loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels = y, logits = scores))  
cross_entropy_loss += reg_strength * tf.reduce_sum(tf.square(W))
```

```
iteration: 700  loss: 1179.57  
iteration: 800  loss: 159.498  
iteration: 900  loss: 23.1094  
iteration: 1000  loss: 4.93028  
iteration: 1100  loss: 2.53405  
iteration: 1200  loss: 2.19926  
iteration: 1300  loss: 2.15682  
iteration: 1400  loss: 2.16726
```

Tensorflow Practice - CIFAR10



Tensorflow Practice - CIFAR10



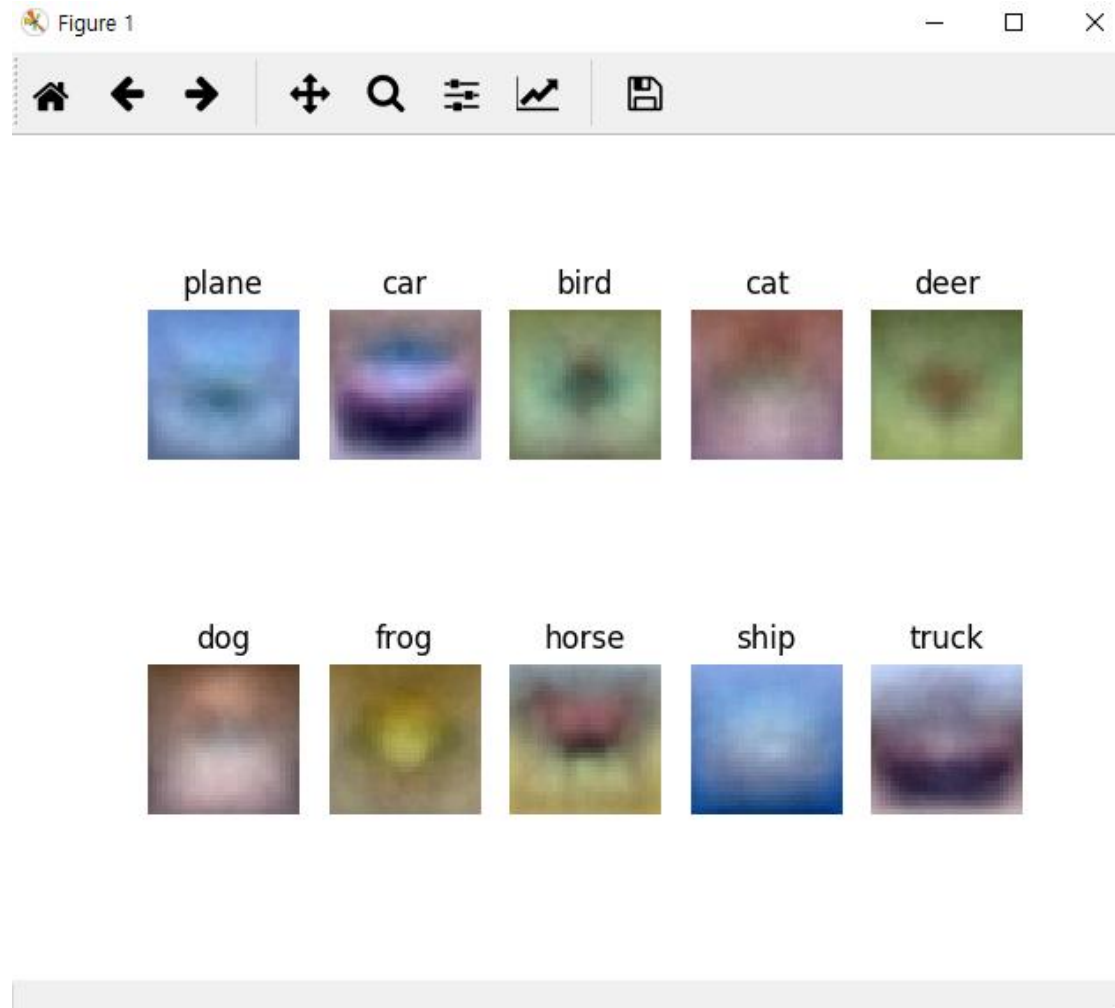
iteration: 1300 loss: 2.15682

iteration: 1400 loss: 2.16726

0.32

32%

Tensorflow Practice - CIFAR10



Next

Image



Array of 32x32x3 numbers
(3072 numbers total)

model is mathematical function!

model

$f(x, W)$

W

parameters or weights

10 numbers giving
class **scores**

Next

Linear Classifier

$$f(x, W) = Wx$$

Image



Array of 32x32x3 numbers
(3072 numbers total)

model

$f(x, W)$

W

parameters or weights

10 numbers giving
class **scores**

Next - Deep Learning

Image



Array of 32x32x3 numbers
(3072 numbers total)


Neural Network
 $f(x, W) =$ 

model

$f(x, W)$

W

parameters or weights



10 numbers giving
class **scores**

Question